

Technical Briefing: Satellite Imagery Segmentation Pipeline for Slum Detection(FPN-VGG)

Executive Summary

This document provides a comprehensive technical overview of a deep learning pipeline designed to detect slum areas through satellite imagery segmentation. Utilizing a Feature Pyramid Network (FPN) with a VGG16 backbone, the model was trained over 30 epochs, ultimately achieving a peak validation F1 score of 0.9482. The methodology incorporates a hybrid loss function (Binary Cross Entropy and Dice Loss) and aggressive data augmentation techniques based on established research. Key performance indicators show a steady improvement in model accuracy, supported by a robust data pipeline that leverages local SSD synchronization and specialized preprocessing for multi-channel GeoTIFF images.

Dataset link: <https://www.arcgis.com/home/item.html?id=e95daeb4085d4afb8ca81d0a1b11dce2>

Research paper link: <https://pdfs.semanticscholar.org/bc49/24f2e32d684372bf7aa10548205b77507a6c.pdf>

VGG16-FPN Slum Segmentation — Project Report

1. Objective

Develop a deep learning model to **automatically identify and segment slum regions** from satellite imagery using semantic segmentation. The model takes a satellite image as input and produces a pixel-level binary mask indicating slum vs. non-slum areas.

2. Dataset

Property	Detail
Source	Satellite GeoTIFF imagery (Google Drive)

Property	Detail
Image format	.tif (3-band RGB, read via rasterio)
Labels	Binary masks — 1 = slum, 0 = non-slum
Split	80% Train / 10% Validation / 10% Test
Split method	Stratified random split (train_test_split, random_state=42)

Only image–label pairs with matching filenames were retained to ensure data integrity.

3. Data Preprocessing & Augmentation

Preprocessing

- Images normalized to 8-bit [0, 255] range, then scaled to [0, 1] as float tensors.
- All images and masks resized to **256 × 256** pixels.

Training Augmentations

To improve generalization and reduce overfitting, the following augmentations were applied to the training set using **Albumentations**:

Augmentation	Parameter
Horizontal Flip	p = 0.5
Vertical Flip	p = 0.5
Random 90° Rotation	p = 0.5

Augmentation	Parameter
Affine (scale, rotate, translate)	scale $\pm 10\%$, rotate $\pm 45^\circ$, translate $\pm 6.25\%$, p = 0.5
Random Brightness/Contrast	p = 0.2

Validation and test sets received **only resizing** (no augmentation) to ensure fair evaluation.

4. Model Architecture

Architecture : Feature Pyramid Network (FPN)

Encoder : VGG-16 (pretrained on ImageNet)

Input : 3-channel RGB, 256×256

Output : 1-channel sigmoid map, 256×256

Library : segmentation_models_pytorch (smp)

Why VGG16-FPN?

- **VGG-16 encoder** provides robust, well-understood feature extraction from ImageNet pre-training, enabling strong transfer learning even with limited satellite imagery data.
- **FPN decoder** aggregates multi-scale features from different encoder stages, allowing the model to capture both fine-grained boundary details and broader contextual information — critical for segmenting irregular slum boundaries.

5. Training Configuration

Hyperparameter	Value
Loss Function	Hybrid Loss = BCELoss + DiceLoss

Hyperparameter	Value
Optimizer	Adam, lr = 1e-4
LR Scheduler	ReduceLROnPlateau (monitor Val F1, factor 0.5, patience 2)
Batch Size	8
Gradient Accumulation	4 steps (effective batch size = 32)
Epochs	30
Model Selection	Best checkpoint saved by highest Validation F1

Hybrid Loss Rationale

- **BCE Loss** ensures accurate per-pixel probability estimation.
- **Dice Loss** directly optimizes spatial overlap, counteracting class imbalance (slum pixels are typically a minority).
- Combining both yields a loss that balances pixel-level accuracy with region-level overlap quality.

6. Training Process

Satellite Image VGG-16 Encoder FPN Decoder Sigmoid Output Hybrid Loss BCE + Dice Adam Optimizer + LR Scheduler

- **Gradient accumulation** (4 steps) allowed an effective batch size of 32 while fitting within GPU memory constraints.
- **ReduceLROnPlateau** scheduler halved the learning rate when validation F1 plateaued for 2 consecutive epochs, enabling finer convergence.

- The **best model checkpoint** was saved based on the highest validation F1 score across all 30 epochs.

Final Training Summary

Metric	Value
Total Epochs	30
Train Loss (final epoch, Epoch 30)	0.7496
Train Loss (best model epoch, Epoch 29)	0.7569
Best Val F1	0.9188 (Epoch 27 & 29)
Training samples	233 batches × 8 = ~1,864 images
Time per epoch	~60 seconds (training) + ~5 seconds (validation)

Epoch-by-Epoch Training Progression

Epoch	Train Loss	Val F1	LR	Best?
1	1.1007	0.7851	1e-4	✓
2	0.8892	0.8469	1e-4	✓
3	0.8594	0.8409	1e-4	

Epoch	Train Loss	Val F1	LR	Best?
4	0.8476	0.8404	1e-4	
5	0.8337	0.8793	1e-4	✓
6	0.8265	0.8609	1e-4	
7	0.8153	0.8794	1e-4	✓
8	0.8065	0.8759	1e-4	
9	0.8009	0.8923	5e-5	✓
10	0.7961	0.8948	5e-5	✓
11	0.7919	0.8840	5e-5	
12	0.7862	0.9009	5e-5	✓
13	0.7881	0.9007	5e-5	
14	0.7814	0.9022	5e-5	✓
15	0.7816	0.9030	5e-5	✓
16	0.7759	0.9000	5e-5	

Epoch	Train Loss	Val F1	LR	Best?
17	0.7783	0.8984	5e-5	
18	0.7779	0.9064	5e-5	✓
19	0.7688	0.9014	5e-5	
20	0.7777	0.9084	5e-5	✓
21	0.7676	0.9107	5e-5	✓
22	0.7700	0.8998	5e-5	
23	0.7671	0.9103	5e-5	
24	0.7671	0.9107	5e-5	
25	0.7563	0.9151	2.5e-5	✓
26	0.7580	0.9162	2.5e-5	✓
27	0.7552	0.9188	2.5e-5	✓
28	0.7557	0.9158	2.5e-5	
29	0.7569	0.9188	2.5e-5	✓

Epoch	Train Loss	Val F1	LR	Best?
30	0.7496	0.9180	2.5e-5	

Key Training Observations

- **Loss dropped rapidly** in early epochs (1.10 → 0.83 in the first 5 epochs), then settled into a slow, steady decline.
- **LR schedule fired twice:** 1e-4 → 5e-5 (around epoch 9) → 2.5e-5 (around epoch 25). Each LR reduction gave a noticeable Val F1 boost.
- **Best model saved at Epoch 29** with Val F1 = 0.9188. The model continued to improve gradually throughout all 30 epochs — no signs of overfitting.
- **Val F1 plateaued** around 0.91–0.92 in the final epochs, suggesting the model has nearly reached its capacity with this architecture and data.

7. Understanding the Loss vs. F1 Gap — Why Loss Looks High but the Model is Actually Great

At first glance, a training loss of **0.7569** might seem concerning — shouldn't a good model have a loss close to zero? But in this case, the high-looking loss number is **completely normal and expected**. Here's why, explained in simple terms.

7.1 We're Adding Two Losses Together

Our loss function is: **Loss = BCE + Dice**. That means the final number is the **sum of two separate penalties**. Think of it like a score on two different exams:

- **BCE Loss** might be around ~0.35 (out of a possible 0)
- **Dice Loss** might be around ~0.40 (out of a possible 0)
- **Combined** = 0.35 + 0.40 = **0.75**

If we had used just BCE alone, the loss would look like ~0.35 — which sounds perfectly fine for a 91% F1 model. We're just stacking two numbers, so the total looks bigger than it really is.

7.2 Loss Measures Confidence, F1 Measures Correctness

This is the most important distinction. Think of it with a simple analogy:

Imagine a student taking a True/False exam. The teacher asks the student: **"How sure are you about each answer?"**

- The student answers 91 out of 100 questions correctly → **F1 $\approx$ 0.91** (great!)
- But the student says "I'm about 75% sure" on most answers instead of "I'm 99% sure" → **Loss stays high** because the student isn't fully confident

The student is **getting the right answers**, but isn't **maximally confident** about them. The loss penalizes that lack of confidence, even though the final answers are correct.

In our model's case:

- When the model predicts a slum pixel with **0.85 confidence** (which is correct after thresholding at 0.5), the F1 score counts it as a win. But the loss function says "you should have been 1.0 confident, not 0.85" and still applies a penalty.
- When the model predicts a non-slum pixel with **0.15 confidence** (also correct), the loss again says "you should have been 0.0, not 0.15" another small penalty.

These small penalties across millions of pixels add up to a loss of ~ 0.76 , even though the model is getting the vast majority of pixels right.

7.3 A Concrete Pixel-by-Pixel Example

Pixel	Truth	Model Prediction	Correct? (F1)	BCE Loss for this pixel
Pixel A (slum)	1.0	0.85	Yes (> 0.5)	$-\log(0.85) = 0.16$
Pixel B (non-slum)	0.0	0.10	Yes (< 0.5)	$-\log(0.90) = 0.11$
Pixel C (slum)	1.0	0.60	Yes (> 0.5)	$-\log(0.60) = 0.51$

Pixel	Truth	Model Prediction	Correct? (F1)	BCE Loss for this pixel
Pixel D (slum)	1.0	0.35	✗ No (< 0.5)	$-\log(0.35) = 1.05$

Notice how Pixels A, B, and C are all **correct** (boosting F1), but they still contribute loss values of 0.16, 0.11, and 0.51. Only Pixel D is actually wrong. The loss accumulates even from correct predictions because they're not perfectly 0 or 1.

7.4 What Would It Take to Get Loss Near Zero?

For the loss to drop to near zero, the model would need to predict **0.99+** for every slum pixel and **0.01-** for every non-slum pixel. That level of extreme confidence usually means the model has **memorized the training data** (overfitting), which is actually **bad** — it would fail on new images.

A moderately confident model (predicting 0.7–0.9 for slums, 0.1–0.3 for non-slums) is healthier and generalizes better to new satellite images it has never seen before.

7.5 Bottom Line

Question	Answer
Is loss = 0.76 bad?	No. It's two losses added together. Each individually is ~0.35–0.40, which is normal.
Why isn't loss near 0?	Because loss penalizes lack of full confidence , not just wrong answers. A correct 0.7 prediction still gets penalized.
Is the model average?	Not at all. F1 = 0.91 and IoU = 0.84 are strong results for segmentation.
What should I trust?	F1 and IoU are your real performance indicators. Loss is just the training signal that guides learning — not a final grade.

8. Test Set Results

The best checkpoint was loaded and evaluated on the held-out 10% test set (data the model never saw during training):

Metric	Score	What It Means (Plain English)
F1 Score	0.9141	Out of all the slum/non-slum decisions the model makes, 91% are correct when balancing both types of mistakes.
IoU (Jaccard)	0.8419	If you overlay the predicted mask on the real mask, they overlap by 84%. Think of it as "how well do the shapes match."
Precision	0.9117	When the model says "this pixel is slum," it's right 91% of the time (very few false alarms).
Recall	0.9166	The model successfully finds 92% of all actual slum pixels (misses very few).
Accuracy	0.9692	97% of all pixels (both slum and non-slum) are classified correctly.

What These Numbers Tell Us

- **Precision $\approx$ Recall** — The model is well-balanced. It doesn't trade finding more slums (recall) for making more false alarms (precision), or vice versa. This balance is ideal.
 - **High IoU (84%)** — The predicted slum boundaries closely match the real boundaries. An IoU above 0.80 is considered a strong result in semantic segmentation.
 - **Accuracy of 97%** — Sounds very high, but this is partly because most of each satellite image is non-slum (background). The model easily gets those right, inflating the number. That's why **F1 and IoU are more reliable metrics** than accuracy for this task.
-

9. Visualization Pipeline

The notebook includes two visualization components:

1. **Test set predictions** — Side-by-side display of: Satellite Image | Ground Truth Mask | Predicted Mask.
 2. **Inference on new images** — The trained model can be applied to completely unseen GeoTIFF images for real-world slum mapping.
-

10. Key Strengths

-  **Transfer learning** from ImageNet pre-trained VGG-16 accelerates convergence and improves generalization
-  **FPN multi-scale fusion** captures both local texture and global context
-  **Heavy augmentation** improves robustness to orientation, scale, and lighting variations
-  **Hybrid loss** directly addresses class imbalance common in satellite segmentation
-  **Balanced precision/recall** — the model does not sacrifice one for the other

11. Potential Improvements

Area	Suggestion
Encoder upgrade	Try ResNet-50, EfficientNet-B4, or a larger backbone for richer features
Post-processing	Apply CRF or morphological ops to clean mask boundaries
Input resolution	Train at 512×512 for finer spatial detail
Loss weighting	Tune BCE:Dice ratio (e.g., 0.5×BCE + 1.0×Dice) to prioritize overlap

Area	Suggestion
Extended training	Use cosine annealing LR schedule with warm restarts for longer training
Test-time augmentation	Average predictions across flips/rotations for more robust inference

12. Conclusion

The VGG16-FPN model demonstrates **strong and reliable semantic segmentation** of slum regions from satellite imagery, achieving an **F1 of 0.9141** and an **IoU of 0.8419** on the held-out test set. The balanced precision and recall indicate a model that is both accurate and comprehensive in its detections. The hybrid BCE + Dice loss effectively handled class imbalance, and the training pipeline — with augmentation, gradient accumulation, and adaptive learning rate scheduling — ensured robust convergence within 30 epochs.

Input Image




Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image




Predicted Slum Mask



Input Image



Predicted Slum Mask




Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image




Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Testing on images other than training dataset:structure of this images(Image width:

600 pixels

Image height: 600 pixels

Number of bands (channels): 4

Image CRS (Coordinate Reference System): EPSG:32643)

Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask




Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask




Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image




Predicted Slum Mask




Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask




Tiff images i downloaded from sam:


From this notebook : https://colab.research.google.com/github/opengeos/segment-geospatial/blob/main/docs/examples/automatic_mask_generator.ipynb#scrollTo=GaUuELD07k9o



Input Image





Predicted Slum Mask




Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask



Input Image



Predicted Slum Mask

